

JAYANT MOHITE

Certified Cloud Data Platform Architect | Big Data Consultant |
Technical Trainer

CONTACT

-  +91-9545-831-218
 hello@jayantmohite.com
 Baner, Pune, Maharashtra, India
 www.jayantmohite.com

EDUCATION

Pursuing
Mumbai university

- Master of Business Management

2008 - 2012
pune University

- Bachelor of Engineering
 - Computer Engineering

CERTIFICATIONS

- Microsoft Certified Azure Solutions Architect Expert
 - Databricks Certified Data Engineer Professional
 - Microsoft Certified Power BI Data Analyst Associate
 - Microsoft Certified Azure Administrator
 - Salesforce Certified Administrator (SCA)
 - Salesforce Certified Advanced Administrator (SCAA)
 - Salesforce Certified Sales Cloud Consultant

PROFILE SUMMARY

Results-driven Azure Solutions Architect & Technical Manager with 13+ years of experience designing and delivering enterprise-scale data platforms across **Azure, Databricks, Microsoft Fabric, and Power BI**. Proven expertise in building end-to-end cloud data ecosystems unifying ingestion, processing, governance, machine learning, & analytics. Strong background in **Hadoop, Spark, Kafka, MongoDB, Cassandra, and Neo4j**, evolved into modern lakehouse architectures leveraging **Delta Lake, Synapse, and Fabric OneLake**. Skilled in **FinOps, CI/CD, data governance, & real-time streaming**, with a track record of driving cost optimization, scalability, and actionable insights for global organizations.

WORK EXPERIENCE

SKAN AI LABS

2020 - Present

Data Platform Manager / Solutions Architect

Transformed traditional data ecosystem into a modern, cloud-native analytics platform built on **Azure**, **Databricks**, **Microsoft Fabric**, and **Power BI**, unifying data, AI, and governance to deliver real-time business intelligence and measurable cost efficiency.

- Partnered with CXOs, product owners, and client stakeholders to define modernization roadmaps, cloud strategies, and reference architectures aligned with business goals.
 - Designed end-to-end architectures integrating **Data Factory**, **Synapse**, **Event Hub**, **Kafka**, **Airflow**, **Blob Storage**, **ADLS**, **OneLake**, and **Delta Lake** for unified batch and streaming workloads.
 - Engineered optimized **PySpark** and Structured Streaming pipelines via Databricks Workflows, **Auto Loader**, and **REST APIs**, achieving a 40% increase in throughput and reduced latency.
 - Led Delta Lake infrastructure design using **Delta Engine**, **Z-Ordering**, and **Data Skipping**, improving query and storage performance; orchestrated the Hive Metastore → Unity Catalog migration with ABAC and ACLs for robust data governance.
 - Led architectural assessments for external client engagements, recommending Hadoop-to-Databricks, SQL-to-Fabric, and on-prem to cloud landing zone migration strategies
 - Managed and modernized **MongoDB** infrastructure, overseeing migration to MongoDB Atlas with Global Clusters, Multi-Region Writes, and Automated Sharding; optimized schema design and indexing for a 30% improvement in query execution times.

TRAINING & PUBLICATIONS

- Best-Seller Author - MongoDB Administration course on Udemy (4.3+ stars, thousands of learners)
- Published Author - Technical content for HP Micro Focus Big Data, O'Reilly, Packt Publications
- Featured in Higher Education Magazine for contributions to skill-based technical education
- Conducted 300+ corporate, academic and retail trainings on Databricks, Big Data, Spark, Power BI, NoSQL Databases and Salesforce
- Certified Salesforce Trainer for 12+ courses
- Founded Wisdom Sprouts in 2016 - An Academic and Retail Training Platform

BIG DATA PROCESSING

- Apache Airflow
- Apache Hadoop
- Apache Nifi
- Apache Pig
- Apache Spark
- Apache Sqoop
- Cloudera Hadoop
- Databricks
- Data Factory
- Data Marts
- MongoDB
- MS Fabric
- One Lake
- Snowflake
- StarRocks
- Synapse Analytics
- Unity Catalog

WORK EXPERIENCE

- Built secure multi-region infrastructure with VNets, private endpoints, and managed identities; executed FinOps and performance-tuning strategies (auto-scaling, caching, parallelization, telemetry) via Azure Monitor and Log Analytics, reducing compute costs by 25%.
- Established data governance, lineage, and monitoring frameworks using Unity Catalog, Purview, Fabric Workspaces, and Log Analytics, ensuring 99.9% uptime for critical operations.
- Architected semantic and visualization layers with Power BI, Fabric DirectLake, and Synapse SQL, delivering governed dashboards and KPI reporting for leadership teams.
- Integrated Azure Machine Learning, Snowflake, PostgreSQL, MongoDB, and Scikit-Learn for predictive analytics and cross-platform data exchange.
- Promoted automated CI/CD and release management practices via Azure DevOps, enabling faster, scalable, and cost-effective pipeline deployments.
- Led the feature engineering team using Agile methodologies, managing sprint planning, reviews, and delivery of high-impact data and analytics features aligned with business objectives.

PALANNTECH INDIA

2017 - 2020

Senior Big Data Consultant

Directed large-scale Big Data platform design and delivery initiatives for enterprise clients, leading cross-functional teams of data engineers and DevOps specialists to modernize data ecosystems across Hadoop, NoSQL, and cloud environments (Azure, AWS).

Key Achievements & Contributions

- Managed agile data-engineering and DevOps teams delivering new analytics capabilities and advanced feature-engineering pipelines under aggressive timelines.
- Designed and deployed end-to-end Hadoop clusters using HDFS, YARN, Hive, Pig, Sqoop, and Spark, enabling distributed analytics on multi-terabyte workloads.
- Built ETL and ELT workflows with Apache Airflow, Oozie, and Kafka, orchestrating ingestion from RDBMS, APIs, and file systems at enterprise scale.
- Architected and administered Cassandra, MongoDB, and Neo4j environments for high-volume, schema-flexible, and relationship-intensive use cases.
- Achieved a 20% reduction in cloud infrastructure spend through improved resource utilization and cross-cloud cost governance (Azure + AWS).
- Implemented partitioning and schema-design strategies in Hive and HBase for optimized query performance and scalability.
- Developed Power BI, Tableau, and Superset dashboards to provide real-time insights and visual analytics to business stakeholders.
- Configured cluster security, capacity planning, and performance tuning, ensuring high availability and efficient workload distribution.
- Provided architecture consulting and migration roadmaps for transitioning on-prem Hadoop workloads to cloud-native data platforms on Azure and AWS.

CLOUD INFRASTRUCTURE

- Azure Infrastructure Management
- Azure Cost Management
- Azure DevOps
- Azure Monitor
- AWS
- Docker
- Elastic Search
- FinOps
- Grafana
- Kibana
- Kubernetes
- Log Analytics
- Prometheus
- Terraform

EVENT STREAMING

- Apache Flink
- Apache Flume
- Apache Kafka
- Event Hub
- RabbitMQ
- Sqoop

DATA VISUALISATION

- Chart.js
- Power BI
- SuperSet
- Tableau

MANAGEMENT

- Agile Delivery Management
- Change Management
- Customer Engagement
- Product Management
- Risk Management
- Strategic Planning & Execution

WORK EXPERIENCE

FREELANCER

2014 - 2017

Big Data Consultant

Delivered end-to-end Big Data and NoSQL solutions for clients across finance, e-commerce, and analytics domains, leveraging the Apache Hadoop ecosystem and modern distributed data technologies.

Key Achievements & Responsibilities

- Built and deployed data platforms using **Apache Hadoop** (**Hortonworks, Cloudera**), **Spark**, **Kafka**, **Cassandra**, **MongoDB**, and **Neo4j**, enabling high-volume analytics and real-time processing.
- Designed custom ETL frameworks with **Apache NiFi**, **Airflow**, and **Pig** to automate ingestion, transformation, and loading workflows, improving reliability and reducing manual intervention.
- Implemented **Kafka + Spark Streaming** pipelines for financial and e-commerce clients, delivering low-latency, high-throughput event-driven analytics.
- Developed scalable data models and **replication/sharding** strategies in **MongoDB**, **Cassandra**, and **Neo4j**, optimizing query performance and fault tolerance.
- Built Hadoop data lakes supporting batch and near real-time workloads with **Hive**, **Pig**, **Sqoop**, and **Spark**, consolidating structured and semi-structured data sources.
- Conducted hands-on client training sessions on Big Data stacks, empowering internal teams to manage and scale their own data infrastructure.
- Created shell/Python-based monitoring and data-flow scripts using **REST APIs** for pipeline observability and operational control.

DB XENTO SYSTEMS

2012 - 2014

Software Engineer

- Administered and optimized **PostgreSQL** databases, improving performance through query tuning and indexing.
- Automated data validation and reconciliation processes using **Python** scripts, reducing manual QA effort and improving data reliability.
- Developed **Power BI** prototypes for internal performance monitoring, combining **SQL** and **Python** outputs into visual insights for engineering managers.
- Monitored and fine-tuned application and database servers for high availability, fault tolerance, and scalability.
- Built reusable ETL utilities in **Python** to extract and transform data from **REST APIs**, **CSVs**, and transactional databases.
- Wrote stored procedures, functions, and triggers to support business workflows and reporting use cases.